

Gender Classification Of Names

Using Machine Learning To Detect/Predict Gender of Individuals From their Names

- Sklearn
- Pandas
- Text Extraction

```
In [1]: import pandas as pd
import numpy as np
import seaborn as sns
```

EDA

```
In [2]: df = pd.read_csv('data/names_dataset.csv')
```

```
In [3]: df.head()
```

```
Out[3]:
```

	index	name	sex
0	0	Mary	F
1	1	Anna	F
2	2	Emma	F
3	3	Elizabeth	F
4	4	Minnie	F

```
In [4]: df.shape
```

```
Out[4]: (95025, 3)
```

```
In [5]: df.drop('index', axis=1, inplace=True)
```


```
In [6]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 95025 entries, 0 to 95024
Data columns (total 2 columns):
#   Column   Non-Null Count  Dtype
---  -
0   name     95025 non-null  object
1   sex      95025 non-null  object
dtypes: object(2)
memory usage: 1.5+ MB
```

```
In [7]: sex = df['sex'].value_counts()
sex
```

```
Out[7]: F      60600
M      34425
Name: sex, dtype: int64
```

```
In [8]: sns.barplot(x=[sex.index[0],sex.index[1]], y=[sex[0],sex[1]]);
```



Encode sex into numbers, Female:0 & Male:1

```
In [9]: # df['sex'].replace({'F':0, 'M':1}, inplace=True)
df['sex'] = df['sex'].apply(lambda x: 0 if x=='F' else 1)
```

```
In [10]: df.head()
```

```
Out[10]:
```

	name	sex
0	Mary	0
1	Anna	0
2	Emma	0
3	Elizabeth	0
4	Minnie	0

```
In [11]: from sklearn.feature_extraction.text import CountVectorizer
```

```
In [12]: cv = CountVectorizer()
```

```
In [13]: X, y = df['name'], df['sex']
```

```
In [14]: df_X = cv.fit_transform(X)
```

In [15]: `print(df_X)`

```
(0, 59607)    1
(1, 5972)     1
(2, 27397)    1
(3, 26638)    1
(4, 62314)    1
(5, 58433)    1
(6, 34639)    1
(7, 3360)     1
(8, 10909)    1
(9, 75946)    1
(10, 6206)    1
(11, 17319)   1
(12, 26692)   1
(13, 29870)   1
(14, 18080)   1
(15, 59475)   1
(16, 52670)   1
(17, 65449)   1
(18, 31970)   1
(19, 14486)   1
(20, 60068)   1
(21, 56459)   1
(22, 10959)   1
(23, 40401)   1
(24, 31305)   1
:            :
(95000, 90488)    1
(95001, 91505)    1
(95002, 91535)    1
(95003, 91826)    1
(95004, 91833)    1
(95005, 91920)    1
(95006, 91999)    1
(95007, 92081)    1
(95008, 92626)    1
(95009, 92629)    1
(95010, 92651)    1
(95011, 92840)    1
(95012, 92908)    1
(95013, 93144)    1
(95014, 93237)    1
(95015, 93402)    1
(95016, 93420)    1
(95017, 93841)    1
(95018, 93862)    1
(95019, 93914)    1
(95020, 94102)    1
(95021, 94463)    1
(95022, 94578)    1
(95023, 94915)    1
(95024, 95019)    1
```

```
In [16]: cv.get_feature_names()
```

```
Out[16]: ['aabab',
           'aabba',
           'aabid',
           'aabriella',
           'aada',
           'aadam',
           'aadan',
           'aadarsh',
           'aaden',
           'aadesh',
           'aadhav',
           'aadhaven',
           'aadi',
           'aadhi',
           'aadhira',
           'aadhvika',
           'aadhya',
           'aadhyan',
           'aadi',
           'adian',
           'aadil',
           'aadin',
           'aadish',
           'aadison',
           'aadit',
           'aadith',
           'aadithya',
           'aaditri',
           'aaditya',
           'aadiv',
           'aadon',
           'aadrian',
           'aadrika',
           'aadrit',
           'aadvik',
           'aadvika',
           'aadya',
           'aadyan',
           'aafia',
           'aafreen',
           'aagam',
           'aage',
           'aagot',
           'aahaan',
           'aahan',
           'aahana',
           'aahil',
           'aahir',
           'aahliyah',
           'aahna',
           'aahron',
           'aaidan',
           'aaiden',
           'aaidyn',
           'aaila',
           'aailiyah',
           'aailyah',
           'aaima',
           'aaira',
           'aaiah',
           'aairah',
           'aisha',
           'aishah',
           'aiyana',
           'aiza',
           'aja',
           'ajah',
           'ajaylah',
           'ajon',
           'aakanksha',
           'aakarsh',
           'aakash',
           'akeem',
           'akilah',
           'akira',
           'akiyah',
           'akriti',
           'ala',
           'alaiya',
```

```
In [17]: df.sort_values('name').head()
```

```
Out[17]:
```

	name	sex
83170	Aaban	1
88709	Aabha	0
75731	Aabid	1
84129	Aabriella	0
94209	Aada	0

```
In [18]: from sklearn.model_selection import train_test_split
X_train, X_test, y_train, y_test = train_test_split(df_X, y,
                                                    train_size=0.75,
                                                    random_state=0)
```

Model Building

```
In [19]: # To measure time of the training process for each algorithm
from time import time
```

Naive Bayes

```
In [20]: from sklearn.naive_bayes import MultinomialNB
nb_clf = MultinomialNB()
tic = time()
nb_clf.fit(X_train, y_train)
toc = time()
```

```
In [21]: # Train set Accuracy
print('Train set Accuracy :', nb_clf.score(X_train, y_train)*100, '%')
# Test set Accuracy
print('Test set Accuracy :', nb_clf.score(X_test, y_test)*100, '%')
# Time of training
print('Time :', (toc-tic)*1000, 'ms')
```

```
Train set Accuracy : 100.0 %
Test set Accuracy : 63.67807383087089 %
Time : 130.2802562713623 ms
```

Logistic Regression

```
In [22]: from sklearn.linear_model import LogisticRegression
lr_clf = LogisticRegression()
tic = time()
lr_clf.fit(X_train, y_train)
toc = time()
```

```
C:\ProgramData\Anaconda3\lib\site-packages\sklearn\linear_model\logistic.py:432: FutureWarning: Default solver will be changed to 'lbfgs' in 0.22. Specify a solver to silence this warning.
FutureWarning)
```

```
In [23]: # Train set Accuracy
print('Train set Accuracy :', lr_clf.score(X_train, y_train)*100, '%')
# Test set Accuracy
print('Test set Accuracy :', lr_clf.score(X_test, y_test)*100, '%')
# Time of training
print('Time :', (toc-tic)*1000, 'ms')
```

```
Train set Accuracy : 63.80423191334119 %
Test set Accuracy : 63.67807383087089 %
Time : 491.89019203186035 ms
```

Note: I tried **Decision Tree & Random Forest** here but they need a lot of time

Sample Prediction

```
In [24]: # Sample 1
sample_name = ['Mary']
vec = cv.transform(sample_name).toarray()
```

```
In [25]: vec
```

```
Out[25]: array([[0, 0, 0, ..., 0, 0, 0]], dtype=int64)
```

```
In [26]: # Female is 0, Male is 1
nb_clf.predict(vec)[0]
```

```
Out[26]: 0
```

```
In [27]: # Sample 2
sample_name2 = ['Omar']
pred2 = nb_clf.predict(cv.transform(sample_name2))
```

```
In [28]: pred2[0]
```

```
Out[28]: 1
```

```
In [29]: # Sample 3
sample_name3 = ['Rania', 'Aya', 'Ahmed', 'Abd', 'Khaled', 'Khan', 'Dal']
pred3 = nb_clf.predict(cv.transform(sample_name3))
pred3
```

```
Out[29]: array([0, 0, 1, 0, 1, 1, 1], dtype=int64)
```

Gender prediction function

```
In [30]: def genderpredictor(name):
         vect = cv.transform([name]).toarray()
         if nb_clf.predict(vect) == 1:
             print('Male')
         else:
             print('Female')
```

```
In [31]: genderpredictor('hanya')
```

```
Female
```

Using a custom function for feature analysis

By analogy, most female names ends in 'A' or 'E' or has the sound of 'A'

```
In [32]: def features(name):
         name = name.lower()
         return {
             'first-letter': name[0], # First letter
             'first2-letters': name[0:2], # First 2 letters
             'first3-letters': name[0:3], # First 3 letters
             'last-letter': name[-1],
             'last2-letters': name[-2:],
             'last3-letters': name[-3:]
         }
```

Example from docs of `vectorize()` function

```
def myfunc(a, b):
    "Return a-b if a>b, otherwise return a+b"
    if a > b:
        return a - b
    else:
        return a + b

vfunc = np.vectorize(myfunc)
vfunc([1, 2, 3, 4], 2)
==> array([3, 4, 1, 2])
```

like

```
np.array([1, 2, 3]) - 1
==> array([0, 1, 2])
```

```
In [33]: # Vectorize the features function
features = np.vectorize(features)
print(features(['Rania', 'Aya', 'Ahmed', 'Abd', 'Khaled', 'Dal']))

[{'first-letter': 'r', 'first2-letters': 'ra', 'first3-letters': 'ran', 'last-letter': 'a', 'last2-letters': 'ia', 'last3-letters': 'nia'}
 {'first-letter': 'a', 'first2-letters': 'ay', 'first3-letters': 'aya', 'last-letter': 'a', 'last2-letters': 'ya', 'last3-letters': 'aya'}
 {'first-letter': 'a', 'first2-letters': 'ah', 'first3-letters': 'ahm', 'last-letter': 'd', 'last2-letters': 'ed', 'last3-letters': 'med'}
 {'first-letter': 'a', 'first2-letters': 'ab', 'first3-letters': 'abd', 'last-letter': 'd', 'last2-letters': 'bd', 'last3-letters': 'abd'}
 {'first-letter': 'k', 'first2-letters': 'kh', 'first3-letters': 'kha', 'last-letter': 'd', 'last2-letters': 'ed', 'last3-letters': 'led'}
 {'first-letter': 'd', 'first2-letters': 'da', 'first3-letters': 'dal', 'last-letter': 'l', 'last2-letters': 'al', 'last3-letters': 'dal'}]
```

```
In [34]: # Extract the features of the dataset
names_features = features(X)
```

```
In [35]: from sklearn.feature_extraction import DictVectorizer
```

what DictVectorizer does is

{'length': 1, 'width': 0, 'height': 2} => {'height': 2, 'length': 1, 'width': 0} => [2, 1, 0]

{'length': 0, 'width': 1, 'height': 1} => {'height': 1, 'length': 0, 'width': 1} => [1, 0, 1]

{'length': 3, 'width': 2, 'height': 1} => {'height': 1, 'length': 3, 'width': 2} => [1, 3, 2]

```
In [36]: v = DictVectorizer(sparse=False)
d = [{'length': 1, 'width': 0, 'height': 2},
     {'length': 0, 'width': 1, 'height': 1},
     {'length': 3, 'width': 2, 'height': 1}]
v.fit_transform(d)
```

```
Out[36]: array([[2., 1., 0.],
                [1., 0., 1.],
                [1., 3., 2.]])
```

```
In [37]: v.get_feature_names()
```

```
Out[37]: ['height', 'length', 'width']
```



```
In [38]: corpus = features(['Amr', 'Omar'])
dv = DictVectorizer(sparse=False)
transformed = dv.fit_transform(corpus)
print('Without sparse:\n', transformed)
print('\nWith sparse:\n', DictVectorizer(sparse=True).fit_transform(corpus))
```

Without sparse:

```
[[1. 0. 1. 0. 1. 0. 1. 0. 1. 1. 0.]
 [0. 1. 0. 1. 0. 1. 1. 1. 0. 0. 1.]]
```

With sparse:

```
(0, 0)      1.0
(0, 2)      1.0
(0, 4)      1.0
(0, 6)      1.0
(0, 8)      1.0
(0, 9)      1.0
(1, 1)      1.0
(1, 3)      1.0
(1, 5)      1.0
(1, 6)      1.0
(1, 7)      1.0
(1, 10)     1.0
```

```
In [39]: dv.get_feature_names()
```

```
Out[39]: ['first-letter=a',
'first-letter=o',
'first2-letters=am',
'first2-letters=om',
'first3-letters=amr',
'first3-letters=oma',
'last-letter=r',
'last2-letters=ar',
'last2-letters=mr',
'last3-letters=amr',
'last3-letters=mar']
```

check the *features* of each name in `corpus` with the `DictVectorizer`'s `get_feature_names` list, if `True` then `1` else `0`.

```
In [40]: X_train2, X_test2, y_train2, y_test2 = train_test_split(names_features, y, train_size=0.75, random_state=0)
```

```
In [41]: X_train2
```

```
Out[41]: array([{'first-letter': 'r', 'first2-letters': 'ru', 'first3-letters': 'rud', 'last-letter': 'a', 'last
2-letters': 'na', 'last3-letters': 'ina'},
 {'first-letter': 'i', 'first2-letters': 'il', 'first3-letters': 'ili', 'last-letter': 'y', 'last
2-letters': 'ny', 'last3-letters': 'nny'},
 {'first-letter': 'b', 'first2-letters': 'bl', 'first3-letters': 'bla', 'last-letter': 'e', 'last
2-letters': 'ie', 'last3-letters': 'lie'},
 ...,
 {'first-letter': 's', 'first2-letters': 'sh', 'first3-letters': 'sha', 'last-letter': 'h', 'last
2-letters': 'ah', 'last3-letters': 'nah'},
 {'first-letter': 'm', 'first2-letters': 'ma', 'first3-letters': 'mar', 'last-letter': 'i', 'last
2-letters': 'ni', 'last3-letters': 'eni'},
 {'first-letter': 's', 'first2-letters': 'se', 'first3-letters': 'sev', 'last-letter': 'n', 'last
2-letters': 'en', 'last3-letters': 'ren'}],
 dtype=object)
```

```
In [42]: # Prepare train/test data
dv = DictVectorizer() # here, sparse can't be False, cuz fitting
X_train2 = dv.fit_transform(X_train2)
X_test2 = dv.transform(X_test2)
```

Model Building

Naive Bayes

```
In [43]: from sklearn.naive_bayes import MultinomialNB
nb_clf = MultinomialNB()
tic = time()
nb_clf.fit(X_train2, y_train2)
toc = time()
nb_time = (toc-tic)*1000
```

```
In [44]: # Train set Accuracy
print('Train set Accuracy :', nb_clf.score(X_test2, y_test2)*100, '%')
# Test set Accuracy
print('Test set Accuracy :', nb_clf.score(X_test2, y_test2)*100, '%')
# Time of training
print('Time :', nb_time, 'ms')
```

```
Train set Accuracy : 85.48217367512733 %
Test set Accuracy : 85.48217367512733 %
Time : 16.989946365356445 ms
```

Logistic Regression

```
In [45]: from sklearn.linear_model import LogisticRegression
lr_clf = LogisticRegression()
tic = time()
lr_clf.fit(X_train2, y_train2)
toc = time()
lr_time = (toc-tic)*1000
```

```
C:\ProgramData\Anaconda3\lib\site-packages\sklearn\linear_model\logistic.py:432: FutureWarning: Default solver will be changed to 'lbfgs' in 0.22. Specify a solver to silence this warning.
  FutureWarning)
```

```
In [46]: # Train set Accuracy
print('Train set Accuracy :', lr_clf.score(X_train2, y_train2)*100, '%')
# Test set Accuracy
print('Test set Accuracy :', lr_clf.score(X_test2, y_test2)*100, '%')
# Time of training
print('Time :', lr_time, 'ms')
```

```
Train set Accuracy : 89.62227086490431 %
Test set Accuracy : 87.46895651807888 %
Time : 1056.7469596862793 ms
```

Decision Tree

```
In [47]: from sklearn.tree import DecisionTreeClassifier
dt_clf = DecisionTreeClassifier()
tic = time()
dt_clf.fit(X_train2, y_train2)
toc = time()
dt_time = (toc-tic)*1000
```

```
In [48]: # Train set Accuracy
print('Train set Accuracy :', dt_clf.score(X_train2, y_train2)*100, '%')
# Test set Accuracy
print('Test set Accuracy :', dt_clf.score(X_test2, y_test2)*100, '%')
# Time of training
print('Time :', dt_time, 'ms')
```

```
Train set Accuracy : 98.79469046416344 %
Test set Accuracy : 86.26931009807636 %
Time : 9620.091438293457 ms
```

Random Forest

```
In [49]: from sklearn.ensemble import RandomForestClassifier
rf_clf = RandomForestClassifier()
tic = time()
rf_clf.fit(X_train2, y_train2)
toc = time()
rf_time = (toc-tic)*1000
```

C:\ProgramData\Anaconda3\lib\site-packages\sklearn\ensemble\forest.py:245: FutureWarning: The default value of n_estimators will change from 10 in version 0.20 to 100 in 0.22.
"10 in version 0.20 to 100 in 0.22.", FutureWarning)

```
In [50]: # Train set Accuracy
print('Train set Accuracy :', rf_clf.score(X_train2, y_train2)*100, '%')
# Test set Accuracy
print('Test set Accuracy :', rf_clf.score(X_test2, y_test2)*100, '%')
# Time of training
print('Time :', rf_time, 'ms')
```

Train set Accuracy : 98.04259976427008 %
Test set Accuracy : 87.15746937744665 %
Time : 19449.76496696472 ms

Sample Prediction

```
In [51]: nb_clf.predict(dv.transform(features(['Hania', 'Haila', 'Abd', 'Mostafa'])))
```

Out[51]: array([0, 0, 1, 0], dtype=int64)

```
In [52]: lr_clf.predict(dv.transform(features(['Hania', 'Haila', 'Abd', 'Mostafa'])))
```

Out[52]: array([0, 0, 1, 1], dtype=int64)

```
In [53]: dt_clf.predict(dv.transform(features(['Hania', 'Haila', 'Abd', 'Mostafa'])))
```

Out[53]: array([0, 0, 1, 0], dtype=int64)

```
In [54]: rf_clf.predict(dv.transform(features(['Hania', 'Haila', 'Abd', 'Mostafa'])))
```

Out[54]: array([0, 0, 1, 1], dtype=int64)

Great progress for *test accuracy* of about **24%** from **63%** to **87%**, specifically for **Logistic Regression & Random Forest**.

But let's visualize a comparison for the 4 models *scores & times*.

Model Evaluation

Scoring Time

```
In [55]: tic = time()
nb_clf_preds = nb_clf.predict(X_train2)
nb_test_time = (time()-tic)*1000
print('NB test time:', nb_test_time, 'ms')

tic = time()
lr_clf_preds = lr_clf.predict(X_train2)
lr_test_time = (time()-tic)*1000
print('LR test time:', lr_test_time, 'ms')

tic = time()
dt_clf_preds = dt_clf.predict(X_train2)
dt_test_time = (time()-tic)*1000
print('DT test time:', dt_test_time, 'ms')


tic = time()
rf_clf_preds = rf_clf.predict(X_train2)
rf_test_time = (time()-tic)*1000
print('RF test time:', rf_test_time, 'ms')
```

```
NB test time: 11.0015869140625 ms
LR test time: 3.0078887939453125 ms
DT test time: 60.225486755371094 ms
RF test time: 474.90596771240234 ms
```

Visualise Socres (only)


```
In [56]: import matplotlib.pyplot as plt
```

```
In [57]: algorithms = ['NB', 'LR', 'DT', 'RF']
scores = np.array([nb_clf.score(X_test2, y_test2), lr_clf.score(X_test2, y_test2),
                    dt_clf.score(X_test2, y_test2), rf_clf.score(X_test2, y_test2)]) * 100
train_times = np.array([nb_time, lr_time, dt_time, rf_time])
test_times = np.array([nb_test_time, lr_test_time, dt_test_time, rf_test_time])
plt.bar(x=algorithms, height=scores);
```



It's little bit difficult to notice the differences between the scores, so we will set the `log` parameter to `True` to normalize the scores

```
In [58]: plt.bar(x=algorithms, height=scores, log=True);
```



Ok, now we can see that **LR** has the best *score* then **RF**, but let's visualize the *time* also by making the times control the bars' widths

Visualise Socres & Times

```
In [59]: # Normalize the times for ease of visualization
train_times = np.log(train_times)
test_times = np.log(test_times)

# put the data into a DataFrame for ease of use
data = pd.DataFrame({
    'score': scores,
    'train_time': train_times,
    'test_time': test_times
}, index=algorithms)
data
```

Out [59]:

	score	train_time	test_time
NB	85.482174	2.832622	2.398040
LR	87.468957	6.962951	1.101238
DT	86.269310	9.171609	4.098096
RF	87.157469	9.875590	6.163117

Visualize Training Time

```


In [60]: ## set figize
_, ax = plt.subplots(figsize=(15,8))

## Draw Bars
ax = sns.barplot(x=data.index, y='score', data=data, log=True, ax=ax)

## Normalize time since bar width is [0,1]
widths = np.array(data.train_time)/10

## Set widths of Bars
# y-axis: represent accuracy of the algorithm
# bar_width: represent time of the training
for bar, newwidth in zip(ax.patches, widths):
    x = bar.get_x()
    width = bar.get_width()
    centre = x + width/2.
    bar.set_x(centre - newwidth/2.)
    bar.set_width(newwidth)

```




Visualize Testing Time

```
In [61]: ## set figize
_, ax = plt.subplots(figsize=(15,8))

## Draw Bars
ax = sns.barplot(x=data.index, y='score', data=data, log=True, ax=ax)

## Normalize time since bar width is [0,1]
widths = np.array(data.test_time)/10

## Set widths of Bars
# y-axis: represent accuracy of the algorithm
# bar_width: represent time of the training
for bar, newwidth in zip(ax.patches, widths):
    x = bar.get_x()
    width = bar.get_width()
    centre = x + width/2.
    bar.set_x(centre - newwidth/2.)
    bar.set_width(newwidth)
```



We can see that **LR** achieved both the best *accuracy* & *time* (test).

Now let's save our model for future use.

Saving Our Model

We can save our model by several methods, ex:

- joblib from sklearn.externals
- pickle

then `.dump(model, destination)` to **save** the model, or `.load(model, destination)` to **load** a model.

You can read more about when to use each one of them from [here \(https://stackoverflow.com/questions/12615525/what-are-the-different-use-cases-of-joblib-versus-pickle\)](https://stackoverflow.com/questions/12615525/what-are-the-different-use-cases-of-joblib-versus-pickle)

joblib is usually significantly faster on large numpy arrays because it has a special handling for the array buffers of the numpy datastructure. To find about the implementation details you can have a look at the [source code](#). It can also compress that data on the fly while pickling using zlib or lz4.

joblib also makes it possible to memory map the data buffer of an uncompressed joblib-pickled numpy array when loading it which makes it possible to share memory between processes.

Note that if you don't pickle large numpy arrays, then regular pickle can be significantly faster, especially on large collections of small python objects (e.g. a large dict of str objects) because the pickle module of the standard library is implemented in C while joblib is pure python.

Note that once PEP 574 (Pickle protocol 5) is merged (hopefully for Python 3.8), it will be much more efficient to pickle large numpy arrays using the standard library.

joblib might still be useful to load objects that have nested numpy arrays in memory mapped mode with `mmap_mode="r"` though.

Save Logistic Regression Model

```
In [62]: from sklearn.externals import joblib    # import pickle
# save the model
joblib.dump(lr_clf, 'models/logregmodel.pkl')    # pickle.dump(_,_)

# load the model
model = joblib.load('models/logregmodel.pkl')    # pickle.load()
```

C:\ProgramData\Anaconda3\lib\site-packages\sklearn\externals\joblib__init__.py:15: DeprecationWarning: sklearn.externals.joblib is deprecated in 0.21 and will be removed in 0.23. Please import this functionality directly from joblib, which can be installed with: pip install joblib. If this warning is raised when loading pickled models, you may need to re-serialize those models with scikit-learn 0.21+. warnings.warn(msg, category=DeprecationWarning)

```
In [64]: model.predict(dv.transform(features(['jack'])))
```

```
Out[64]: array([1], dtype=int64)
```

Make a Pipeline

To make it easy for prediction to the application, we will merge the `DictVectorizer` with the `LogisticRegression` classifier into a pipeline.

```
In [65]: from sklearn.pipeline import Pipeline

vec_clf = Pipeline([('vectorizer', dv), ('lr', lr_clf)])
# vec_clf.fit(X_train2, y_train2)
joblib.dump(vec_clf, 'models/vectorizer_and_lr.pkl')
```

```
Out[65]: ['models/vectorizer_and_lr.pkl']
```

```
In [66]: model = joblib.load('models/vectorizer_and_lr.pkl')
```

```
In [67]: model.predict(features(['jack']))
```

```
Out[67]: array([1], dtype=int64)
```

```
In [ ]:
```